

Naming the Intelligence Layer & Agent Family

Recommendation memo · Forwardable to Kathrina, Chris, Neil

Situation

The UK trademark application for **Comply365 CoAnalyst** is facing opposition, and the Board has asked how our AI capability should be branded as the market shifts from generic Co-prefixed assistants (Copilot, CoAnalyst) to named AI personas with distinct agent families (Salesforce **Einstein** + **Agentforce**, SAP **Joule**, ServiceNow **Now Assist**, Optimizely **Opal**). Internally we have already moved the narrative to the **Intelligence Layer**; externally the legacy CoAnalyst name still appears in legal filings and a small number of artefacts.

Recommendation

- Retire **CoAnalyst** from all customer-facing surfaces. Withdraw or narrow the UK TM rather than fight a multi-jurisdiction opposition.
- Standardise on **Intelligence Layer** as the architectural term — it already describes what the capability is, not a sub-brand.
- Land a proper **persona name** for the Intelligence Layer (the Tier-1 AI brand). Run a 60–90 day naming sprint with formal TM clearance in UK / EU / US / Madrid.
- Keep **DTOP** as the four capabilities of the layer — **not** as four agents. DTOP is the operating model, not a SKU shelf.
- Introduce a **Comply365 Agents** family. Each agent is a **Job-To-Be-Done**, slotted under the DTOP stage where the job lives, owned by one application.

Why now

Every major operational SaaS has now landed a named AI persona plus an agent family in the last 18 months. Customers, analysts and regulators are starting to ask the question in those terms. Holding a contested Co-prefix slows the conversation and dates the brand before it ships at scale.

The ask

Board approval for: (a) retiring CoAnalyst externally, (b) the three-tier naming convention in §2, (c) the Capabilities-vs-Agents JTBD model in §3–§4, and (d) a 60–90 day naming sprint to land the Intelligence Layer persona name (shortlist in §5).

The Three-Tier Naming Convention

How every Comply365 brand asset hangs together


Tier 1	Platform	Comply365 — the Operational Performance Platform. Master brand. Branded-house architecture.
Tier 2	Applications	SafetyManager365 · ContentManager365 · TrainingManager365 . BrandNumber format. No spaces. Each app is a self-contained product surface.
Tier 3a	Intelligence Layer (persona)	[PersonaName] — one ownable, trademarkable AI brand sitting across all three applications. The Tier-1 AI brand (like Einstein, Joule).
Tier 3b	Capabilities (DTOP)	The four stages of the layer: Detect (Insights) · Trigger (Intelligence & Recommendations) · Orchestrate (Automation) · Prove (Outcomes & Audit).
Tier 3c	Agents (JTBD)	Comply365 Agents family. Each agent is a Job-To-Be-Done, slotted under one DTOP stage, owned by one application. Naming pattern: [Job] Agent .

Naming rules (non-negotiable)

- No **Co-** prefixes anywhere. The category has moved on.
- No sub-brand collisions with applications (no Comply365 ContentX, etc.).
- The persona name is never used as a verb (we don't "[Persona] it").
- DTOP stages are **capabilities**, never agents. They describe how the layer works.
- Every agent is named after the **job**, mapped to one DTOP stage, with one owning application.
- Agent family name: **Comply365 Agents**. Singular agent name: **[Job] Agent**.

The Architecture, At A Glance

One platform · three apps · one intelligence layer · four capabilities · many agents



The diagram is the contract. Every new AI capability we ship must map cleanly onto one DTOP stage, be owned by one application, and — if it does a human job — be named as a Comply365 Agent.

Capabilities vs Agents — the JTBD Model

The rule that stops everything becoming "an agent"

	Capabilities	Agents
What it is	A stage of the pipeline. <i>What the layer does.</i>	A named worker doing a Job-To-Be-Done a human would otherwise do.
How many	Exactly four. Detect · Trigger · Orchestrate · Prove.	Many. A growing catalogue, versioned over time.
Named for	The operating model.	The job. Pattern: [Job] Agent .
Sold as	Part of the Intelligence Layer — not sold individually.	Packaged inside applications; some cross-app under Platform.
The test	"Is this a stage every job flows through?" → Yes.	"Would a customer staff this role today?" → Yes.

Why DTOP must stay as capabilities

- DTOP is our **operating model**, not a SKU shelf. Rebranding the four stages as four agents creates a packaging trap — customers will ask to price, license and turn each on/off individually.
- Every competitor can claim a Detect/Trigger/Orchestrate/Prove agent. The defensibility is in **domain-trained intelligence powering named JTBD agents**, not in the stage names.
- It dilutes the persona name we're about to land. If everything is an agent, the Tier-1 AI brand has no room to breathe.
- It confuses the message customers have already learned: Insights, Intelligence & Recommendations, Automation, Outcomes & Audit are **product modules** — keep them as such.

Why agents must be JTBD, not features

- Buyers don't fund features; they fund **jobs that get done**. Naming agents after the job makes the value visible on the org chart.
- JTBD agents survive product reorganisations — the Authoring Agent does the same job whether it lives in ContentManager365 or moves to a new module.
- It maps cleanly to ROI: each agent has a baseline human cost, a target automation rate and a measurable outcome.

Comply365 Agents — v1 Catalogue

Every agent is a JTBD, slotted under one DTOP stage, owned by one application

DTOP stage	JTBD agent	The job (in one line)	Owning app
■ Detect	Safety Signal Agent	Surface emerging risk patterns across reports before they become events.	SafetyManager365
■ Detect	Fleet Risk Agent	Monitor operational data streams and flag fleet-level anomalies.	SafetyManager365
■ Detect	Reg Change Agent	Watch regulator feeds and flag changes that affect live manuals.	ContentManager365
■ Trigger	Root Cause Agent	Classify and cite the cause / root cause of an event at L4–5 accuracy.	SafetyManager365
■ Trigger	Compliance Review Agent	Assess a change against the regulatory framework and recommend actions.	ContentManager365
■ Trigger	Competency Gap Agent	Match incident patterns to training records and surface skill gaps.	TrainingManager365
■ Orchestrate	Authoring Agent	Draft the procedure / manual update from the recommended action.	ContentManager365
■ Orchestrate	Training Update Agent	Generate the matching training delta and route it to affected crew.	TrainingManager365
■ Orchestrate	Workflow Agent	Route the right task to the right person or system with SLAs.	Platform
■ Prove	Audit Trail Agent	Assemble end-to-end evidence: signal → decision → change → outcome.	Platform
■ Prove	Outcomes Agent	Quantify value delivered and feed it back into the next Detect cycle.	Platform

Locked naming pattern. [Job] Agent — short, descriptive, role-based. Family name: **Comply365 Agents**. New agents must (a) name a real human job, (b) map to exactly one DTOP stage, (c) declare an owning application, (d) declare baseline human cost and target automation rate before launch.

Naming Shortlist — the Intelligence Layer Persona

12 candidates screened on ownability, sector fit, pronounceability, intelligence connotation, and initial TM-class-9/42 desk risk (not legal clearance)

Candidate	Pool	Rationale	Initial risk
Vector	Navigation	Direction + intelligence. Strong aviation resonance. Short, ownable.	● Green
Atlas	Navigation	Authority, knowledge, maps. Memorable but widely used — clearance risk in software class.	● Amber
Beacon	Navigation	Signal + guidance. On-message for Detect → Prove. Some crowding in safety tech.	● Amber
Compass	Navigation	Direction setting. Strong metaphor but heavily used in finance/HR SaaS.	● Red
Orbit	Navigation	Continuous loop — fits DTOP cycle. Distinctive in aviation context.	● Green
Horizon	Navigation	Forward-looking, predictive. Already used by several aviation programmes.	● Red
Pulse	Signal	Living, continuous signal. Strong fit for Detect. Crowded in healthtech.	● Amber
Sentry	Signal	Watchful, protective. Strong for safety. Risk: feels reactive, not predictive.	● Amber
Halo	Signal	Surrounds the platform. Distinctive. Existing brand collisions (gaming, automotive).	● Red
Lumen	Signal	Light, clarity, intelligence. Short, ownable, soft. Some collision in lighting/SaaS.	● Amber
Crew	Crew	Operator metaphor — your people, augmented. Generic word, hard to TM standalone.	● Red
Cadre	Crew	Trained, trusted team. Distinctive, ownable, aviation-credible.	● Green

Top 3 recommendation

- **Vector** — strongest aviation resonance, cleanest sound, easiest sell internally. Implies direction and intelligence in one word. **Recommended primary.**
- **Orbit** — best metaphor for the DTOP cycle (Detect → Trigger → Orchestrate → Prove → Detect). Distinctive in our space.
- **Cadre** — best fit if we want the persona to feel like a trusted operator alongside the crew, not a tool.

60–90 day path to launch

- **Weeks 1–2:** ELT alignment on top 3; brief external naming/IP firm (Stobbs or Mishcon).
- **Weeks 3–6:** Formal TM clearance in UK / EU / US / Madrid Protocol, classes 9 and 42. Domain and social handle sweep.
- **Weeks 7–10:** Customer-tested narrative — 2 live customer councils, 1 analyst preview (Gartner / IDC).
- **Weeks 11–12:** Lock name. Update decks, site, playbooks. Withdraw or narrow the CoAnalyst UK TM in parallel.

Appendix

A. Competitor AI brand architectures

Vendor	Persona (Tier 1 AI brand)	Agent family	Capabilities exposed as
Salesforce	Einstein	Agentforce (named agents)	Pipeline stages, not branded
SAP	Joule	Joule Agents (role-based)	Capabilities behind the persona
ServiceNow	Now Assist	AI Agents (per workflow)	Workflow stages, not branded
Optimizely	Opal	Opal Agents (per job)	Implicit; not surfaced as SKUs
Microsoft	Copilot	Copilot Agents (Studio-built)	Capabilities behind the persona

Pattern: every Tier-1 vendor has now landed (a) one persona, (b) a named agent family, (c) capabilities hidden behind the persona. Our proposed architecture mirrors this pattern with the DTOP operating model as our defensible difference.

B. Risk register — staying with "CoAnalyst"

Risk	Severity	Mitigation if we stay
Multi-jurisdiction TM opposition (UK opened; EU/US likely)	● High	18+ months of legal cost; outcome uncertain.
Co- prefix dating the brand vs Einstein/Joule/Opal	● High	Re-platform narrative in 18–24 months anyway.
Persona confusion with Copilot / CoAnalyst (generic SaaS pattern)	● Med	Heavier brand spend to claim airspace.
Internal divergence — decks already say "Intelligence Layer"	● Med	Forced rework of recently shipped surfaces.
Agent family naming gets blocked by Tier-3 persona	● Med	Cannot land "Comply365 Agents" cleanly under a contested name.

Out of scope for this brief

Formal legal TM clearance (next step with external counsel). Updating any in-product copy, decks or playbooks (once a persona name is approved, a follow-on rollout plan covers those surfaces). Picking the final persona name — this brief shortlists and recommends; the Board/ELT decides.